
Assignment #4

1. **[noSQL]** (40) This set of questions get you to engage discussion related with noSQL systems.
 - a. Explain the concept of noSQL
 - b. Describe the 4 types of noSQL systems; for each category, give an example noSQL system (try to give different system from what we introduced in the class).
 - c. Pick one type of noSQL system you give in b, give a real-world application that can make best use of the system, and an application that may not be well-suited for the system. Explain your reasons.

2. **[CAP and ACID]** (20) This set of questions are related to data consistency
 - a. What is CAP Theory? Consider a toy example of a cluster that contains two servers S1 and S2. Use the example cluster to explain the CAP theory. You can simply use the proof I explained in class.

 - b. Consider the relation Accounts(acctNo, balance), and two SQL statements that conduct a request “transfer \$200 from account 123 to 456”. In a DBMS, this usually includes two steps:
 - i. Add \$200 to account 456:
`UPDATE Accounts SET balance=balance+200 WHERE acctNo = 456`
 - ii. Subtract \$200 from account 123:
`UPDATE Accounts SET balance=balance - 200 WHERE acctNo=123`

Use this example and necessary scenarios to show when Atomicity, Consistency, Isolation and Durability can be violated.

3. **[Column Store]** (20) We take a closer look at Column Store
 - a. Explain the major features of column stores in terms of data storage and storage key.
 - b. We introduced three techniques to optimize column-oriented databases: compression, late materialization and block iteration. Please explain how they work.

4. **[newSQL]** (20)
 - a. Describe the definition of NewSQL systems and design principles.
 - b. For in-memory DBMS, describe the set-associative cache for block placement and LRU block replacement policy.